

Specification for the FIRRTL Language:
Version 0.1.3
PRE-RELEASE VERSION - DO NOT
DISTRIBUTE

Patrick S. Li Adam M. Izraelevitz
psli@eecs.berkeley.edu adamiz@eecs.berkeley.edu

Jonathan Bachrach
jrb@eecs.berkeley.edu

January 21, 2016

Contents

1	Introduction	5
1.1	Background	5
1.2	Design Philosophy	6
2	Acknowledgements	7
3	Circuits and Modules	8
3.1	Circuits	8
3.2	Modules	8
3.3	Externally Defined Modules	9
4	Types	9
4.1	Ground Types	9
4.1.1	Integer Types	9
4.1.2	Clock Type	10
4.2	Vector Types	10

4.3	Bundle Types	10
4.4	Passive Types	11
4.5	Type Equivalence	12
4.6	Weak Type Equivalence	12
4.6.1	Oriented Types	12
4.6.2	Conversion to Oriented Types	13
4.6.3	Oriented Type Equivalence	13
5	Statements	13
5.1	Connect Statements	13
5.1.1	The Connection Algorithm	14
5.2	Partial Connect Statements	15
5.2.1	The Partial Connection Algorithm	16
5.3	Statement Groups	16
5.3.1	Last Connect Semantics	17
5.4	The Empty Statement	18
5.5	Wires	18
5.6	Registers	19
5.7	Invalidate Statements	19
5.7.1	The Invalidate Algorithm	20
5.8	Nodes	20
5.9	The Conditional Statement	21
5.9.1	Syntactic Shorthands	21
5.9.2	Nested Declarations	23
5.9.3	Initialization Coverage	24
5.9.4	Scoping	24
5.9.5	Conditional Last Connect Semantics	24
5.10	Memories	26
5.10.1	Read Ports	28
5.10.2	Write Ports	28
5.10.3	Readwrite Ports	29
5.10.4	Read Under Write Behaviour	29
5.11	Instances	29
5.12	Stops	30
5.13	Formatted Prints	30
5.13.1	Format Strings	31

6	Expressions	31
6.1	Unsigned Integers	32
6.2	Signed Integers	32
6.3	Unsigned Bits	32
6.4	Signed Bits	33
6.5	References	34
6.6	Subfields	34
6.7	Subindices	34
6.8	Subaccesses	35
6.9	Multiplexors	37
6.10	Conditionally Valids	38
6.11	Primitive Operations	38
7	Primitive Operations	39
7.1	Add Operation	39
7.2	Subtract Operation	39
7.3	Multiply Operation	40
7.4	Divide Operation	40
7.5	Modulus Operation	40
7.6	Comparison Operations	41
7.7	Padding Operations	41
7.8	Interpret As UInt	41
7.9	Interpret As SInt	41
7.10	Shift Left Operation	42
7.11	Shift Right Operation	42
7.12	Dynamic Shift Left Operation	42
7.13	Dynamic Shift Right Operation	42
7.14	Arithmetic Convert to Signed Operation	43
7.15	Negate Operation	43
7.16	Bitwise Complement Operation	43
7.17	Binary Bitwise Operations	44
7.18	Bitwise Reduction Operations	44
7.19	Concatenate Operation	44
7.20	Bit Extraction Operation	45
7.21	Head	45
7.22	Tail	45
8	Genders	45

9	Namespaces	46
9.1	Expansion Algorithm	46
9.2	Prefix Uniqueness	47
10	Details about Syntax	47
11	The Lowered FIRRTL Form	48
12	TODO	49
13	FIRRTL Language Definition	50
13.1	Abstract Syntax Tree	50
13.2	Notation	51

1 Introduction

1.1 Background

The ideas for FIRRTL originated from work on Chisel, a hardware description language embedded in Scala used for writing highly-parameterized circuit design generators. Chisel designers manipulate circuit components using Scala functions, encode their interfaces in Scala types, and use Scala's object-orientation features to write their own circuit libraries. This form of meta-programming enables expressive, reliable and type-safe generators that improve RTL design productivity and robustness.

The computer architecture research at U.C. Berkeley relied critically on Chisel to allow small teams of graduate students to design sophisticated RTL circuits. Over a three year period with under twelve graduate students, the architecture team has taped-out over ten different designs.

Internally, the investment in developing and learning Chisel was rewarded with huge gains in productivity. However, Chisel's external rate of adoption was slow for the following reasons.

1. Learning a functional programming language (Scala) is difficult for RTL designers with limited software-engineering experience.
2. Confounding the previous point, conceptually separating the Chisel HDL from the host language is difficult for new users.
3. The output of Chisel (Verilog) is unreadable and slow to simulate.
4. Writing custom circuit transformers requires intimate knowledge about the internals of the Chisel compiler.
5. Chisel semantics are ill-defined and thus impossible to target from other languages.
6. Error checking is unprincipled due to ill-defined semantics resulting in incomprehensible error messages.

As a consequence, Chisel needed to be redesigned from the ground up to standardize its semantics, modularize its compilation process for robustness, and cleanly separate its front-end, intermediate representation, and back-ends. A well defined intermediate representation (IR) allows the system to be targeted by other host programming languages, making it possible for RTL designers to work within a language they are already comfortable with. A clearly defined IR with a concrete syntax also allows inspection of the output of circuit generators/transformers thus making clear the distinction

between the host language and the constructed circuit. A clearly defined semantics allows users without knowledge of the compiler implementation to write circuit transformers. Examples include optimization of circuits for simulation speed, and automatic insertion of signal activity counters. An additional benefit of a well defined IR is the structural invariants that can be enforced before and after each compilation stage, resulting in a more robust compiler and structured mechanism for error checking.

1.2 Design Philosophy

FIRRTL (Flexible Intermediate Representation for RTL) represents the standardized elaborated circuit that the Chisel DSL produces. FIRRTL represents the circuit immediately after Chisel’s elaboration but before any circuit simplification. It is designed to resemble the Chisel DSL after all meta-programming has executed. Thus a user program that makes little use of meta-programming facilities should look almost identical to the generated FIRRTL.

For this reason, FIRRTL has first-class support for high-level constructs such as vector types, bundle types, conditional statements, partial connects, and modules. These high-level constructs are then gradually removed by a sequence of *lowering* transformations. During each lowering transformation the circuit is rewritten into an equivalent circuit using simpler, lower-level constructs. Eventually the circuit is simplified to its most restricted form, resembling a structured netlist, which allows for easy translation to an output language (e.g. Verilog). This form is given the name *lowered FIRRTL* (LoFIRRTL) and is a strict subset of the full FIRRTL language.

Because the host language is now used solely for its meta-programming facilities, the frontend can be very light-weight, and additional frontends in other languages can target FIRRTL and reuse the majority of the compiler toolchain.

Similar to backends, it is often convenient to write transformers that accept only the restricted LoFIRRTL subset. However, the transformed circuit is allowed to contain any FIRRTL construct, as it can be subsequently lowered again. We intentionally designed LoFIRRTL to be a subset of the full FIRRTL language to provide this feature.

2 Acknowledgements

The FIRRTL language could not have been developed without the help of many of the faculty and students in the ASPIRE lab, and the University of California, Berkeley.

This project originated from discussions with our advisor, Jonathan Bachrach, who indicated the need for a structural redesign of the Chisel system around a well-defined intermediate representation. Patrick Li designed and implemented the first prototype of the FIRRTL language, wrote the initial specification for the language, and presented it to the Chisel group consisting of Scott Beamer, David Biancolin, Christopher Celio, Henry Cook, Palmer Dabbelt, Donggyu Kim, Jack Koenig, Martin Maas, Albert Magyar, Colin Schmidt, Andrew Waterman, Yunsup Lee, Richard Lin, Eric Love, Albert Ou, Stephen Twigg, Jim Lawson, Brian Richards, Krste Asanovic, and John Wawrzynek.

Adam Izraelevitz then led the design and reimplemented FIRRTL, and after many discussions with the Chisel group, refined the design to its present version.

The authors would like to thank the following individuals for their contributions to the FIRRTL project:

1. Andrew Waterman: for his many contributions to the design of FIRRTL's constructs, for his work on Chisel 3.0, and for porting architecture research infrastructure
2. Richard Lin: for improving the Chisel 3.0 code base for release quality
3. Jack Koenig: for implementing the FIRRTL parser in Scala
4. Henry Cook: for porting and cleaning up many aspects of Chisel 3.0, including the testing infrastructure and the parameterization library
5. Stephen Twigg: for his expertise in hardware intermediate representations and for providing many corner cases to consider
6. Palmer Dabbelt, Eric Love, Martin Maas, Christopher Celio, and Scott Beamer: for their feedback on previous drafts of the FIRRTL specification

And finally this project would not have been possible without the continuous feedback and encouragement of our advisor, Jonathan Bachrach, and his leadership on and implementation of Chisel 3.0.

Research is partially funded by DARPA Award Number XXXX, the Center for Future Architectures Research, a member of STARnet, a Semiconductor Research Corporation program sponsored by MARCO and DARPA, and ASPIRE Lab industrial sponsors and affiliates Intel, Google, Nokia, NVIDIA, Oracle, and Samsung. Any opinions, findings, conclusions, or recommendations in this paper are solely those of the authors and do not necessarily reflect the position or the policy of the sponsors.

3 Circuits and Modules

3.1 Circuits

All FIRRTL circuits consist of a list of modules, each representing one hardware block that could be instantiated. The circuit must specify the name of the top-level module.

```
circuit MyTop :  
  module MyTop :  
    ...  
    module MyModule :  
      ...
```

3.2 Modules

Each module has a given name, a list of ports, and a statement representing the circuit connections within the module. A module port is specified by its *direction*, which may be input or output, a name, and the data type for the port.

```
module MyModule :  
  input foo: UInt  
  output bar: UInt  
  bar <= foo
```


Note that a module definition does *not* indicate that the module will be physically present in the final circuit. Refer to the description of the instance statement for details on how to instantiate a module (section ??).

3.3 Externally Defined Modules

Externally defined modules consist of a given name, and a list of ports, whose types and names must match its external definition.

```
module MyExternalModule :  
  input foo: UInt  
  output bar: UInt  
  output baz: SInt
```

4 Types

4.1 Ground Types

All types in FIRRTL are either one of the fundamental ground types or are built up from aggregating other types. There are three ground types in FIRRTL, an unsigned integer type, a signed integer type, and a clock type.

4.1.1 Integer Types

Both unsigned and signed integer types may optionally be given a known positive integer bit width.

```
UInt<10>  
SInt<32>
```

Alternatively, if the bit width is omitted, it will be automatically inferred by FIRRTL's width inferencer, as detailed in section ??.

```
UInt  
SInt
```

4.1.2 Clock Type

The clock type is used to describe wires and ports meant for carrying clock signals. The usage of components with clock types are restricted. Clock signals cannot be used in most primitive operations, and clock signals can only be connected to components that have been declared with the clock type.

The clock type is specified as follows:

`Clock`

4.2 Vector Types

A vector type is used to express an ordered sequence of elements of a given type. The length of the sequence must be non-negative and known. This is akin to array types in the C programming language.

The following example specifies a ten element vector of 16-bit unsigned integers.

`UInt<16>[10]`

The next example specifies a ten element vector of unsigned integers of omitted but identical bit widths.

`UInt[10]`

Note that any type, including other aggregate types, may be used as the element type in the sequence. The following example specifies a twenty element vector, each of which is a ten element vector of 16-bit unsigned integers.

`UInt<16>[10][20]`

4.3 Bundle Types

A bundle type is used to express a collection of nested and named types. All fields in a bundle type must have a given name, and type.

The following is an example of a possible type for representing a complex number. It has two fields, `real`, and `imag`, both 10-bit signed integers.

`{real:SInt<10>, imag:SInt<10>}`

Additionally, a field may optionally be declared with a *flipped* orientation.

```
{word:UInt<32>, valid:UInt<1>, flip ready:UInt<1>}
```

In a connection between elements declared with the same bundle type, the data carried by the flipped fields flow in the opposite direction as the data carried by the non-flipped fields.

As an example, consider a module output port declared with the following type:

```
output mysignal: {word:UInt<32>, valid:UInt<1>, flip ready:UInt<1>}
```

In a connection to the `mysignal` port, the data carried by the `word` and `valid` subfields will flow out of the module, while data carried by the `ready` subfield will flow into the module. More details about how the bundle field orientation affects connections are explained in section ??.

As in the case of vector types, a bundle field may be declared with any type, including other aggregate types.

```
{real: {word:UInt<32>, valid:UInt<1>, flip ready:UInt<1>}  
  imag: {word:UInt<32>, valid:UInt<1>, flip ready:UInt<1>}}
```

When calculating the final direction of data flow, the orientation of a field is applied recursively to all nested types in the field. As an example, consider the following module port declared with a bundle type containing a nested bundle type.

```
output myport: {a: UInt, flip b: {c: UInt, flip d:UInt}}
```

In a connection to `myport`, the `a` subfield flows out of the module. The `c` subfield contained in the `b` subfield flows into the module, and the `d` subfield contained in the `b` subfield flows out of the module.

4.4 Passive Types

Intuitively, a passive type is defined as a type where all leaf elements have a non-flipped orientation. Thus all ground types are passive types. Vector types are passive if their element type is passive. And bundle types are passive if no fields are flipped and if all field types are passive.

4.5 Type Equivalence

The type equivalence relation is used to determine whether a connection between two elements is legal. See section ?? for further details about connect statements.

An unsigned integer type is always equivalent to another unsigned integer type regardless of bit width, and is not equivalent to any other type. Similarly, a signed integer type is always equivalent to another signed integer type regardless of bit width, and is not equivalent to any other type.

Clock types are only equivalent to clock types, and are not equivalent to any other type.

Two vector types are equivalent if they have the same length, and if their element types are equivalent.

Two bundle types are equivalent when they have the same number of fields, and the *i*'th field from each bundle have matching names and orientations, as well as equivalent types. Consequently, `{a:UInt, b:UInt}` is a different type than `{b:UInt, a:UInt}`. `{a: {flip b:UInt}}` is a different type than `{flip a: {b: UInt}}`.

4.6 Weak Type Equivalence

The weak type equivalence relation is used to determine whether a partial connection between two elements is legal. See section ?? for further details about partial connect statements.

Two types are weakly equivalent if their corresponding oriented types are equivalent.

4.6.1 Oriented Types

The weak type equivalence relation requires first a definition of *oriented types*. Intuitively, an oriented type is a type where all orientation information is collated and coupled with the leaf ground types instead of in bundle fields.

An oriented ground type is an orientation coupled with a ground type. An oriented vector type is an ordered sequence of elements of a given oriented type. The length of an oriented vector type must be known and positive. An oriented bundle type is a collection of fields containing a name and an oriented type, but no orientation.

Applying a flip orientation to an oriented type recursively reverses the orientation of every oriented ground type contained within. Applying a non-flip orientation to an oriented type does nothing.

4.6.2 Conversion to Oriented Types

To convert a ground type to an oriented ground type, attach a non-flip orientation to the ground type.

To convert a vector type to an oriented vector type, convert its element type to an oriented type, and retain its length.

To convert a bundle field to an oriented bundle field, first convert its type to an oriented type. Then apply the field orientation to the oriented type, returning a new oriented type. The new oriented type, and the original field's name combine to create the oriented bundle field. To convert a bundle type to an oriented bundle type, convert each bundle field to an oriented bundle field.

4.6.3 Oriented Type Equivalence

Two oriented ground types are equivalent if their orientations and types are equivalent.

Two oriented vector types are equivalent if their element types are equivalent.

Two oriented bundle types are not equivalent if there exists two fields, one from each oriented bundle type, that have identical names but whose oriented types are not equivalent. Otherwise, the oriented bundle types are equivalent.

As stated earlier, two types are weakly equivalent if their corresponding oriented types are equivalent.

5 Statements

Statements are used to instantiate and connect circuit elements together.

5.1 Connect Statements

The connect statement is used to specify a physically wired connection between two circuit elements.

The following example demonstrates connecting a module's input port to its output port, where port `myinput` is connected to port `myoutput`.

```
module MyModule :  
  input myinput: UInt  
  output myoutput: UInt  
  myoutput <= myinput
```

In order for a connection to be legal the following conditions must hold:

1. The types of the left-hand and right-hand side expressions must be equivalent (see section ?? for details).
2. The bit widths of the two expressions must allow for data to always flow from a smaller to equal size or larger bit width.
3. The gender of the left-hand side expression must be female or bi-gender (see section ?? for an explanation of gender).
4. Either the gender of the right-hand side expression is male or bi-gender, or the right-hand side expression has a passive type.

Connect statements from a narrower ground type element to a wider ground type element will have its value automatically sign-extended to the larger bit width.

5.1.1 The Connection Algorithm

Connect statements between ground types cannot be expanded further.

Connect statements between two vector typed elements recursively connects each subelement in the right-hand side expression to the corresponding subelement in the left-hand side expression.

Connect statements between two bundle typed elements connects the *i*'th field of the right-hand side expression and the *i*'th field of the left-hand side expression. If the *i*'th field is not flipped, then the right-hand side field is connected to the left-hand side field. Conversely, if the *i*'th field is flipped, then the left-hand side field is connected to the right-hand side field.

5.2 Partial Connect Statements

The partial connect statement is also used to specify a physically wired connection between two circuit elements. However, it enforces fewer restrictions on the types and widths of the circuit elements it connects.

In order for a partial connect to be legal the following conditions must hold:

1. The types of the left-hand and right-hand side expressions must be weakly equivalent (see section ?? for details).
2. The gender of the left-hand side expression must be female or bi-gender (see section ?? for an explanation of gender).
3. Either the gender of the right-hand side expression is male or bi-gender, or the right-hand side expression has a passive type.

Partial connect statements from a narrower ground type element to a wider ground type element will have its value automatically sign-extended to the larger bit width. Partial connect statements from a wider ground type element to a narrower ground type element will have its value automatically truncated to fit the smaller bit width.

Intuitively, bundle fields with matching names will be connected appropriately, while bundle fields not present in both types will be ignored. Similarly, vectors with mismatched lengths will be connected up to the shorter length, and the remaining elements are ignored.

The following example demonstrates partially connecting a module's input port to its output port, where port `myinput` is connected to port `myoutput`.

```
module MyModule :
  input myinput: {flip a:UInt, b:UInt[2]}
  output myoutput: {flip a:UInt, b:UInt[3], c:UInt}
  myoutput <- myinput
```

The above example is equivalent to the following:

```
module MyModule :
  input myinput: {flip a:UInt, b:UInt[2]}
  output myoutput: {flip a:UInt, b:UInt[3], c:UInt}
  myinput.a <- myoutput.a
  myoutput.b[0] <- myinput.b[0]
  myoutput.b[1] <- myinput.b[1]
```

For details on the syntax and semantics of the subfield and subindex expressions, see section ??.

5.2.1 The Partial Connection Algorithm

A partial connect statement between two ground type elements connects the right-hand side expression to the left-hand side expression. Conversely, a reverse partial connect statement between two ground type elements connects the left-hand side expression to the right-hand side expression.

A partial (or reverse partial) connect statement between two vector typed elements applies a partial (or reverse partial) connect from the first n subelements in the right-hand side expression to the first n corresponding subelements in the left-hand side expression, where n is the length of the smaller vector.

A partial (or reverse partial) connect statement between two bundle typed elements considers any pair of fields, one from the first bundle type and one from the second, with matching names. If the first field in the pair is not flipped, then we apply a partial (or reverse partial) connect from the right-hand side field to the left-hand side field. However, if the first field is flipped, then we apply a reverse partial (or partial) connect from the right-hand side field to the left-hand side field.

5.3 Statement Groups

An ordered sequence of one or more statements can be grouped into a single statement, called a statement group. The following example demonstrates a statement group composed of three connect statements.

```
module MyModule :
  input a : UInt
  input b : UInt
  output myport1 : UInt
  output myport2 : UInt
  myport1 <= a
  myport1 <= b
  myport2 <= a
```


5.3.1 Last Connect Semantics

Ordering of statements is significant in a statement group. Intuitively, statements execute in order, and the effects of later statements take precedence over earlier ones. In the previous example, in the resultant circuit, port `b` will be connected to `myport1`, and port `a` will be connected to `myport2`.

Note that `connect` and `partial connect` statements have equal priority, and later `connect` or `partial connect` statements always take priority over earlier `connect` or `partial connect` statements. For details on the behavior of conditional statements under last connect semantics, see section ??.

In the case where a `connect` to a circuit element with an aggregate type is followed by a `connect` to a subcomponent of that element, only the connection to the subcomponent is overwritten. Connections to the other subcomponents remain unaffected. In the following example, in the resultant circuit, the `c` subcomponent of port `portx` will be connected to the `c` subcomponent of `myport`, and port `porty` will be connected to the `b` subcomponent of `myport`.

```
module MyModule :
  input portx: {b:UInt, c:UInt}
  input porty: UInt
  output myport: {b:UInt, c:UInt}
  myport <= portx
  myport.b <= porty
```

The above circuit can be rewritten as follows.

```
module MyModule :
  input portx: {b:UInt, c:UInt}
  input porty: UInt
  output myport: {b:UInt, c:UInt}
  myport.b <= porty
  myport.c <= portx.c
```

In the case where a connection to a subcomponent of an aggregate circuit element is followed by a connection to the entire circuit element, the later connection overwrites the earlier connections completely.

```
module MyModule :
  input portx: {b:UInt, c:UInt}
```

```
input porty: UInt
output myport: {b:UInt, c:UInt}
myport.b <= porty
myport <= portx
```

The above circuit can be rewritten as follows.

```
module MyModule :
  input portx: {b:UInt, c:UInt}
  input porty: UInt
  output myport: {b:UInt, c:UInt}
  myport <= portx
```

See section ?? for more details about subcomponent expressions.

5.4 The Empty Statement

The empty statement does nothing and is used simply as a placeholder where a statement is expected. It is specified using the `skip` keyword.

The following example:

```
a <= b
skip
c <= d
```

can be equivalently expressed as:

```
a <= b
c <= d
```

The empty statement is most often used as the alternate branch in a conditional statement, or as a convenient placeholder for removed components during transformational passes. See section ?? for details on the conditional statement.

5.5 Wires

A wire is a named combinational circuit element that can be connected to and from using connect and partial connect statements.

The following example demonstrates instantiating a wire with the given name `mywire` and type `UInt`.

```
wire mywire : UInt
```

5.6 Registers

A register is a named stateful circuit element.

The following example demonstrates instantiating a register with the given name `myreg`, type `SInt`, and is driven by the clock signal `myclock`.

```
wire myclock: Clock
reg myreg: SInt, myclock
```

Optionally, for the purposes of circuit initialization, a register can be declared with a reset signal and value. In the following example, `myreg` is assigned the value `myinit` when the signal `myreset` is high.

```
wire myclock: Clock
wire myreset: UInt<1>
wire myinit: SInt
reg myreg: SInt, myclock, myreset, myinit
```

Note that the clock signal for a register must be of type `clock`, the reset signal must be a single bit `UInt`, and the type of initialization value must match the declared type of the register.

5.7 Invalidate Statements

An invalidate statement is used to indicate that a circuit component contains indeterminate values. It is specified as follows:

```
wire w:UInt
w is invalid
```

Invalidate statements can be applied to any circuit element of any type. However, if the circuit element cannot be connected to, then the statement has no effect on the element. The following example demonstrates the effect of invalidating a variety of circuit elements with aggregate types. See section ?? for details on the algorithm for determining what is invalidated.

```
module MyModule :
  input in:{flip a:UInt, b:UInt}
  output out:{flip a:UInt, b:UInt}
```

```

wire w:{flip a:UInt, b:UInt}
in is invalid
out is invalid
w is invalid

```

is equivalent to the following:

```

module MyModule :
  input in:{flip a:UInt, b:UInt}
  output out:{flip a:UInt, b:UInt}
  wire w:{flip a:UInt, b:UInt}
  in.a is invalid
  out.b is invalid
  w.a is invalid
  w.b is invalid

```

For the purposes of simulation, invalidated elements are initialized to random values, and operations involving indeterminate values produce undefined behaviour. This is useful for early detection of errors in simulation.

5.7.1 The Invalidate Algorithm

Invalidating an element with a ground type indicates that the element's value is indetermined if the element is female or bi-gender. Otherwise, the element is unaffected.

Invalidating an element with a vector type recursively invalidates each subelement in the vector.

Invalidating an element with a bundle type recursively invalidates each subelement in the bundle.

5.8 Nodes

A node is simply a named intermediate value in a circuit. The node must be initialized to a value with a passive type and cannot be connected to.

The following example demonstrates instantiating a node with the given name `mynode` initialized with the output of a multiplexor.

```

wire pred: UInt<1>
wire a: SInt
wire b: SInt
node mynode = mux(pred, a, b)

```

5.9 The Conditional Statement

The conditional statement is used to specify a condition under which connections to previously declared components hold. The condition must have a 1-bit unsigned integer type.

In the following example, the wire `x` is connected to the wire `a` only when the `en` signal is high. Otherwise, the wire `x` is connected to the wire `b`.

```
module mymodule :  
  input a: UInt  
  input b: UInt  
  input en: UInt<1>  
  wire x: UInt  
  when en :  
    x <= a  
  else :  
    x <= b
```

5.9.1 Syntactic Shorthands

The `else` branch of a conditional statement may be omitted, in which case a default `else` branch is supplied consisting of the empty statement.

Thus the following example:

```
module mymodule :  
  input a: UInt  
  input b: UInt  
  input en: UInt<1>  
  wire x: UInt  
  when en :  
    x <= a
```

can be equivalently expressed as:

```
module mymodule :  
  input a: UInt  
  input b: UInt  
  input en: UInt<1>  
  wire x: UInt
```

```
when en :  
  x <= a  
else :  
  skip
```

To aid readability of long chains of conditional statements, the colon following the **else** keyword may be omitted if the **else** branch consists of a single conditional statement.

Thus the following example:

```
module mymodule :  
  input a: UInt  
  input b: UInt  
  input c: UInt  
  input d: UInt  
  input c1: UInt<1>  
  input c2: UInt<1>  
  input c3: UInt<1>  
  
  wire x: UInt  
  when c1 :  
    x <= a  
  else :  
    when c2 :  
      x <= b  
    else :  
      when c3 :  
        x <= c  
      else :  
        x <= d
```

can be equivalently expressed as:

```
module mymodule :  
  input a: UInt  
  input b: UInt  
  input c: UInt  
  input d: UInt
```

```
input c1: UInt<1>
input c2: UInt<1>
input c3: UInt<1>
```

```
wire x: UInt
when c1 :
  x <= a
else when c2 :
  x <= b
else when c3 :
  x <= c
else :
  x <= d
```

5.9.2 Nested Declarations

If a component is declared within a conditional statement, connections to the component are unaffected by the condition. In the following example, register `myreg1` is always connected to `a`, and register `myreg2` is always connected to `b`.

```
module mymodule :
  input a: UInt
  input b: UInt
  input en: UInt<1>
  input clk : Clock
  when en :
    reg myreg1 : UInt, clk
    myreg1 <= a
  else :
    reg myreg2 : UInt, clk
    myreg2 <= b
```

Intuitively, a line can be drawn between a connection (or partial connection) to an element and that element's declaration. All conditional statements that are crossed by the line apply to that connection (or partial connection).

5.9.3 Initialization Coverage

Because of the conditional statement, it is possible to syntactically express circuits containing wires that are only partially connected to an expression. In the following example, the wire *w* is connected to *a* when *en* is asserted high, but it is not specified what *w* is connected to when *en* is low.

```
wire en: UInt<1>
wire w: UInt
wire a: UInt
when en :
  w <= a
```

This is an illegal FIRRTL circuit and an error will be thrown during compilation. All wires, memory ports, instance ports, and module ports that can be connected to must be connected to under all conditions.

5.9.4 Scoping

The conditional statement creates a new *scope* each within its *when* and *else* branches. It is an error to refer to any component declared within a branch after the branch has ended.

5.9.5 Conditional Last Connect Semantics

In the case where a connect to a circuit element is followed by a conditional statement containing a connect to the same element, the connection is overwritten only when the condition holds. Intuitively, a multiplexor is generated such that when the condition is low, the multiplexor returns the old value, and otherwise returns the new value. For details about the multiplexor, see section ??.

The following example:

```
wire a: UInt
wire b: UInt
wire c: UInt<1>
wire w: UInt
w <= a
when c :
  w <= b
```


can be rewritten equivalently using a multiplexor as follows:

```
wire a: UInt
wire b: UInt
wire c: UInt<1>
wire w: UInt
w <= mux(c, b, a)
```

In the case where an invalid statement is followed by a conditional statement containing a connect to the invalidated element, the resulting connection to the element can be expressed using a conditionally valid expression. See section ?? for more details about the conditionally valid expression.

```
wire a: UInt
wire c: UInt<1>
wire w: UInt
w is invalid
when c :
  w <= a
```

can be rewritten equivalently as follows:

```
wire a: UInt
wire c: UInt<1>
wire w: UInt
w <= validif(c, a)
```

The behaviour of conditional connections to circuit elements with aggregate types can be modeled by first expanding each connect into individual connect statements on its ground elements (see section ?? and ?? for the connection and partial connection algorithms) and then applying the conditional last connect semantics.

The following example:

```
wire x: {a:UInt, b:UInt}
wire y: {a:UInt, b:UInt}
wire c: UInt<1>
wire w: {a:UInt, b:UInt}
w <= x
when c :
  w <= y
```

can be rewritten equivalently as follows:

```
wire x: {a:UInt, b:UInt}
wire y: {a:UInt, b:UInt}
wire c: UInt<1>
wire w: {a:UInt, b:UInt}
w.a <= mux(c, y.a, x.a)
w.b <= mux(c, y.b, x.b)
```

Similar to the behavior of aggregate types under last connect semantics (see section ??), the conditional connects to a subcomponent of an aggregate element only generates a multiplexors for the subcomponents that are overwritten.

```
wire x: {a:UInt, b:UInt}
wire y: UInt
wire c: UInt<1>
wire w: {a:UInt, b:UInt}
w <= x
when c :
  w.a <= y
```

```
wire x: {a:UInt, b:UInt}
wire y: UInt
wire c: UInt<1>
wire w: {a:UInt, b:UInt}
w.a <= mux(c, y, x.a)
w.b <= x.b
```

5.10 Memories

A memory is an abstract representation of a hardware memory. It is characterized by the following parameters.

1. A passive type representing the type of each element in the memory.
2. A positive integer representing the number of elements in the memory.
3. A variable number of named ports, each being a read port, a write port, or readwrite port.

4. A non-negative integer indicating the read latency, which is the number of cycles after setting the port's read address before the corresponding element's value can be read from the port's data field.
5. A non-negative integer indicating the write latency, which is the number of cycles after setting the port's write address/data before the corresponding element within the memory holds the new value.
6. A read-under-write flag indicating the behaviour of the memory when a memory location is written to while a read to that location is in progress.

The following example demonstrates instantiating a memory containing 256 complex numbers, each with 16-bit signed integer fields for its real and imaginary components. It has two read ports, `r1` and `r2`, and one write port, `w`. It is combinational read (read latency is zero cycles) and has a write latency of one cycle. Finally, its read-under-write behavior is undefined.

```
mem mymem :
  data-type => {real:SInt<16>, imag:SInt<16>}
  depth => 256
  reader => r1
  reader => r2
  writer => w
  read-latency => 0
  write-latency => 1
  read-under-write => undefined
```

In the example above, the type of `mymem` is:

```
{flip r1: {flip data:{real:SInt<16>, imag:SInt<16>},
  addr:UInt<8>,
  en:UInt<1>,
  clk:Clock}
flip r2: {flip data:{real:SInt<16>, imag:SInt<16>},
  addr:UInt<8>,
  en:UInt<1>,
  clk:Clock}
flip w: {data:{real:SInt<16>, imag:SInt<16>},
  mask:{real:UInt<1>, imag:UInt<1>},
```

```

    addr:UInt<8>,
    en:UInt<1>,
    clk:Clock}}

```

The following sections describe how a memory's field types are calculated and the behavior of each type of memory port.

5.10.1 Read Ports

If a memory is declared with element type `T`, has a size less than or equal to 2^N , then the corresponding type of its read ports is:

```
{flip data:T, addr:UInt<N>, en:UInt<1>, clk:Clock}
```

If the `en` field is high, then the element associated with the address in the `addr` field can be retrieved by reading from the `data` field after the appropriate read latency. If the `en` field is low, then the value in the `data` field, after the appropriate read latency, is undefined. The port is driven by the clock signal in the `clk` field.

5.10.2 Write Ports

If a memory is declared with element type `T`, has a size less than or equal to 2^N , then the corresponding type of its write ports is:

```
{data:T, mask:M, addr:UInt<N>, en:UInt<1>, clk:Clock}
```

where `M` is the mask type calculated from the element type `T`. Intuitively, the mask type mirrors the aggregate structure of the element type except with all ground types replaced with a single bit unsigned integer type. The *non-masked portion* of the data value is defined as the set of data value subcomponents where the corresponding mask subelement is high.

If the `en` field is high, then the non-masked portion of the `data` field value is written, after the appropriate write latency, to the location indicated by the `addr` field. If the `en` field is low, then no value is written after the appropriate write latency. The port is driven by the clock signal in the `clk` field.

5.10.3 Readwrite Ports

Finally, the corresponding type of its readwrite ports is:

```
{rmode:UInt<1>, flip rdata:T, data:T, mask:M,
  addr:UInt<N>, en:UInt<1>, clk:Clock}
```

A readwrite port is a single port that, on a given cycle, can be used either as a write or a read port. If the readwrite port is in read mode (the `rmode` field is high), then the `rdata`, `addr`, `en`, and `clk` fields constitute its read port fields, and should be used accordingly. If the readwrite port is not in read mode (the `rmode` field is low), then the `data`, `mask`, `addr`, `en`, and `clk` fields constitute its write port fields, and should be used accordingly.

5.10.4 Read Under Write Behaviour

The read-under-write flag indicates the resultant value held on the `data` field of a read port if the memory location is written to while a read to that location is in progress. The flag may take on three settings: `old`, `new`, and `undefined`.

If the read-under-write flag is set to `old`, then a read port always returns the value existing in the memory on the same cycle that the read was requested. Intuitively, this is modeled as a combinational read from the memory that is then delayed by the appropriate read latency.

If the read-under-write flag is set to `new`, then a read port always returns the value existing in the memory on the same cycle that the read was made available. Intuitively, this is modeled as a combinational read from the memory after delaying the read address by the appropriate read latency.

If the read-under-write flag is set to `undefined`, then the value held by the read port after the appropriate read latency is undefined.

In all cases, if a memory location is written to by more than one port on the same cycle, the stored value is undefined.

5.11 Instances

FIRRTL modules are instantiated with the instance statement. The following example demonstrates creating an instance named `myinstance` of the `MyModule` module within the top level module `Top`.

```
circuit Top :  
  module MyModule :  
    input a: UInt  
    output b: UInt  
    b <= a  
  module Top :  
    inst myinstance of MyModule
```

The resulting instance has a bundle type, where each field represents a port, and shares its name, in the instantiated module. The fields corresponding to input ports are flipped to indicate their data flows in the opposite direction as the output ports. The `myinstance` instance in the example above has type `{flip a:UInt, b:UInt}`.

Modules have the property that instances can always be *inlined* into the parent module without affecting the semantics of the circuit.

To disallow infinitely recursive hardware, modules cannot contain instances of itself, either directly, or indirectly through instances of other modules it instantiates.

5.12 Stops

The stop statement is used to halt simulations of the circuit. Backends are free to generate hardware to stop a running circuit for the purpose of debugging, but this is not required by the FIRRTL specification.

A stop statement requires a clock signal, a halt condition signal that has a single bit unsigned integer type, and an integer exit code.

```
wire clk:Clock  
wire halt:UInt<1>  
stop(clk, halt, 42)
```

5.13 Formatted Prints

The formatted print statement is used to print a formatted string during simulations of the circuit. Backends are free to generate hardware that relays this information to a hardware test harness, but this is not required by the FIRRTL specification.

A printf statement requires a clock signal, a print condition signal, a format string, and a variable list of argument signals. The conditional signal

must be a single bit unsigned integer type, and the argument signals must each have a ground type.

```
wire clk:Clock
wire condition:UInt<1>
wire a:UInt
wire b:UInt
printf(clk, condition, "a in hex: %x, b in decimal:%d.\n", a, b)
```

On each positive clock edge, when the conditional signal is high, the printf statement prints out the format string where argument placeholders are substituted with the value of the appropriate argument.

5.13.1 Format Strings

Format strings support the following argument placeholders:

- `%b` : Prints the argument in binary
- `%d` : Prints the argument in decimal
- `%x` : Prints the argument in hexadecimal
- `%%` : Prints a single `%` character

Format strings support the following escape characters:

- `\n` : New line
- `\t` : Tab
- `\\` : Back slash
- `\"` : Double quote
- `\'` : Single quote

6 Expressions

FIRRTL expressions are used for creating literal unsigned and signed integers, for referring to a declared circuit component, for statically and dynamically accessing a nested element within a component, and for performing primitive operations.

6.1 Unsigned Integers

A literal unsigned integer can be created given a non-negative integer value and an optional positive bit width. The following example creates a 10-bit unsigned integer representing the number 42.

```
UInt<10>(42)
```

Note that it is an error to supply a bit width that is not large enough to fit the given value. If the bit width is omitted, then FIRRTL's width inferencer to infer the minimum number of bits necessary to fit the given value.

```
UInt(42)
```

6.2 Signed Integers

Similar to unsigned integers, a literal signed integer can be created given an integer value and an optional positive bit width. The following example creates a 10-bit signed integer representing the number -42.

```
SInt<10>(-42)
```

Note that it is an error to supply a bit width that is not large enough to fit the given value using two's complement representation. If the bit width is omitted, then FIRRTL's width inferencer to infer the minimum number of bits necessary to fit the given value.

```
SInt(-42)
```

6.3 Unsigned Bits

A literal unsigned integer can alternatively be created given a string representing its bit representation and an optional bit width.

The following radices are supported:

1. `0b` : For representing binary numbers.
2. `0o` : For representing octal numbers.
3. `0x` : For representing hexadecimal numbers.

If a bit width is not given, the number of bits in the bit representation is directly represented by the string. The following examples create a 8-bit integer representing the number 13.

```
UBits("0b00001101")  
UBits("0x0D")
```

If a bit width is given, then the bit representation is truncated to the given bit width. It is an error to supply a bit width that is larger than the number of bits in the bit representation. The following examples create a 7-bit integer representing the number 13.

```
UBits<7>("0b00001101")  
UBits<7>("0o015")  
UBits<7>("0x0D")
```

6.4 Signed Bits

Similar to unsigned integers, a literal signed integer can alternatively be created given a string representing its bit representation and an optional bit width.

If a bit width is not given, the number of bits in the bit representation is directly represented by the string. The following examples create a 8-bit integer representing the number -13.

```
SBits("0b11110011")  
SBits("0xF3")
```

If a bit width is given, then the bit representation is truncated to the given bit width. It is an error to supply a bit width that is larger than the number of bits in the bit representation. The following examples create a 7-bit integer representing the number -13.

```
SBits<7>("0b11110011")  
SBits<7>("0o763")  
SBits<7>("0xF3")
```

6.5 References

A reference is simply a name that refers to a previously declared circuit component. It may refer to a module port, node, wire, register, instance, or memory.

The following example connects a reference expression `in`, referring to the previously declared port `in`, to the reference expression `out`, referring to the previously declared port `out`.

```
module MyModule :  
  input in : UInt  
  output out : UInt  
  out <= in
```

In the rest of the document, for brevity, the names of components will be used to refer to a reference expression to that component. Thus, the above example will be rewritten as “the port `in` is connected to the port `out`”.

6.6 Subfields

The subfield expression refers to a subelement of an expression with a bundle type.

The following example connects the `in` port to the `a` subelement of the `out` port.

```
module MyModule :  
  input in : UInt  
  output out : {a:UInt, b:UInt}  
  out.a <= in
```

6.7 Subindices

The subindex expression statically refers, by index, to a subelement of an expression with a vector type. The index must be a non-negative integer and cannot be equal to or exceed the length of the vector it indexes.

The following example connects the `in` port to the fifth subelement of the `out` port.

```

module MyModule :
  input in : UInt
  output out : UInt[10]
  out[4] <= in

```

6.8 Subaccesses

The subaccess expression dynamically refers to a subelement of an expression with a vector type by a calculated index. The index must be an expression with an unsigned integer type.

The following example connects the `n`'th subelement of the `in` port to the `out` port.

```

module MyModule :
  input in: UInt[3]
  input n: UInt<2>
  output out: UInt
  out <= in[n]

```

A connection from a subaccess expression can be modeled by conditionally connecting from every subelement in the vector, where the condition holds when the dynamic index is equal to the subelement's static index.

```

module MyModule :
  input in: UInt[3]
  input n: UInt<2>
  output out: UInt
  when eq(n, UInt(0)) :
    out <= in[0]
  else when eq(n, UInt(1)) :
    out <= in[1]
  else when eq(n, UInt(2)) :
    out <= in[2]
  else :
    out is invalid

```

The following example connects the `in` port to the `n`'th subelement of the `out` port. All other subelements of the `out` port are connected from the corresponding subelements of the `default` port.

```

module MyModule :
  input in: UInt
  input default: UInt[3]
  input n: UInt<2>
  output out: UInt[3]
  out <= default
  out[n] <= in

```

A connection to a subaccess expression can be modeled by conditionally connecting to every subelement in the vector, where the condition holds when the dynamic index is equal to the subelement's static index.

```

module MyModule :
  input in: UInt
  input default: UInt[3]
  input n: UInt<2>
  output out: UInt[3]
  out <= default
  when eq(n, UInt(0)) :
    out[0] <= in
  else when eq(n, UInt(1)) :
    out[1] <= in
  else when eq(n, UInt(2)) :
    out[2] <= in

```

The following example connects the `in` port to the `m`'th subelement of the `n`'th subelement vector of the `out` port. All other subelements of the `out` port are connected from the corresponding subelements of the `default` port.

```

module MyModule :
  input in: UInt
  input default: UInt[2][2]
  input n: UInt<1>
  input m: UInt<1>
  output out: UInt[2][2]
  out <= default
  out[n][m] <= in

```

A connection to an expression containing multiple nested subaccess expressions can also be modeled by conditionally connecting to every subele-

ment in the expression. However the condition holds only when all dynamic indices are equal to all of the subelement's static indices.

```

module MyModule :
  input in: UInt
  input default: UInt[2][2]
  input n: UInt<1>
  input m: UInt<1>
  output out: UInt[2][2]
  out <= default
  when and(eq(n, UInt(0)), eq(m, UInt(0))) :
    out[0][0] <= in
  else when and(eq(n, UInt(0)), eq(m, UInt(1))) :
    out[0][1] <= in
  else when and(eq(n, UInt(1)), eq(m, UInt(0))) :
    out[1][0] <= in
  else when and(eq(n, UInt(1)), eq(m, UInt(1))) :
    out[1][1] <= in

```

6.9 Multiplexors

A multiplexor outputs one of two input expressions depending on the value of an unsigned single bit selection signal.

The following example connects the result of selecting between the **a** and **b** ports to the **c** port. The **a** port is selected when the **sel** signal is high, otherwise the **b** port is selected.

```

module MyModule :
  input a: UInt
  input b: UInt
  input sel: UInt<1>
  output c: UInt
  c <= mux(sel, a, b)

```

A multiplexor expression is legal only if the following holds.

1. The type of the selection signal is a single bit unsigned integer.
2. The types of the two input expressions are equivalent.
3. The types of the two input expressions are passive.

6.10 Conditionally Valids

A conditionally valid expression is expressed as an input expression guarded with an unsigned single bit valid signal. It outputs the input expression when the valid signal is high, otherwise the result is undefined.

The following example connects the `a` port to the `c` port when the `valid` signal is high. Otherwise, the value of the `c` port is undefined.

```
module MyModule :
  input a: UInt
  input valid: UInt<1>
  output c: UInt
  c <= validif(valid, a)
```

A conditionally valid expression is legal only if the following holds.

1. The type of the valid signal is a single bit unsigned integer.
2. The type of the input expression is passive.

Conditional statements are transformed into multiplexors and conditionally valid expressions during lowering to LoFIRRTL. See section ?? for details about conditional statements and section ?? for details about LoFIRRTL and the lowering transformation.

6.11 Primitive Operations

All fundamental operations on ground types are expressed as a FIRRTL primitive operation. In general, each operation takes some number of argument expressions, along with some number of static integer literal parameters.

The general form of a primitive operation is expressed as follows:

```
op(arg0, arg1, ..., argn, int0, int1, ..., intm)
```

The following examples of primitive operations demonstrate adding two expressions, `a` and `b`, shifting expression `a` left by 3 bits, selecting the fourth bit through and including the seventh bit in the `a` expression, and interpreting the expression `x` as a Clock typed signal.

```
add(a, b)
shl(a, 3)
bits(a, 7, 4)
asClock(x)
```

Section 7 will describe the format and semantics of each primitive operation.

7 Primitive Operations

All primitive operations expression operands must be ground types. Each specific operation can place additional restrictions on the number and types of their arguments.

7.1 Add Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
add	(e1, e2)	()	(UInt, UInt)	UInt	$\max(w_{e1}, w_{e2}) + 1$
			(UInt, SInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$
			(SInt, UInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$
			(SInt, SInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$

The add operation result is the sum of e1 and e2 without loss of precision.

7.2 Subtract Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
sub	(e1, e2)	()	(UInt, UInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$
			(UInt, SInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$
			(SInt, UInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$
			(SInt, SInt)	SInt	$\max(w_{e1}, w_{e2}) + 1$

The subtract operation result is e2 subtracted from e1, without loss of precision.

7.3 Multiply Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
mul	(e1, e2)	()	(UInt, UInt)	UInt	$w_{e1} + w_{e2}$
			(UInt, SInt)	SInt	$w_{e1} + w_{e2}$
			(SInt, UInt)	SInt	$w_{e1} + w_{e2}$
			(SInt, SInt)	SInt	$w_{e1} + w_{e2}$

The multiply operation result is the product of `e1` and `e2`, without loss of precision.

7.4 Divide Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
div	(num, den)	()	(UInt, UInt)	UInt	w_{num}
			(UInt, SInt)	SInt	$w_{num} + 1$
			(SInt, UInt)	SInt	w_{num}
			(SInt, SInt)	SInt	$w_{num} + 1$

The divide operation divides `num` by `den`, truncating the fractional portion of the result. This is equivalent to rounding the result towards zero.

7.5 Modulus Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
mod	(num, den)	()	(UInt, UInt)	UInt	$\min(w_{num}, w_{den})$
			(UInt, SInt)	UInt	$\min(w_{num}, w_{den})$
			(SInt, UInt)	SInt	$\min(w_{num}, w_{den} + 1)$
			(SInt, SInt)	SInt	$\min(w_{num}, w_{den})$

The modulus operation yields the remainder from dividing `num` by `den`, keeping the sign of the numerator. Together with the divide operator, the modulus operator satisfies the relationship below:

```
num = add(mul(den, div(num, den)), mod(num, den))
```


7.6 Comparison Operations

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
lt, leq, gt, geq, eq, neq	(e1, e2)	()	(UInt, UInt)	UInt	1
			(UInt, SInt)	UInt	1
			(SInt, UInt)	UInt	1
			(SInt, SInt)	UInt	1

The comparison operations returns an unsigned 1 bit signal with value one if **e1** is less than (**lt**), less than or equal to (**leq**), greater than (**gt**), or greater than or equal to (**geq**), or equal to (**eq**), or not equal to (**neq**) **e2**. The operation returns a value of zero otherwise.

7.7 Padding Operations

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
pad	(e)	(n)	(UInt)	UInt	$\max(w_e, n)$
			(SInt)	SInt	$\max(w_e, n)$

If **e**'s bit width is smaller than **n**, then the pad operation sign-extends or zero-extends **e** up to the given width **n**. Otherwise, the result is simply **e**.

7.8 Interpret As UInt

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
asUInt	(e)	()	(UInt)	UInt	w_e
			(SInt)	UInt	w_e

The interpret as UInt operation reinterprets **e**'s bits as an unsigned integer.

7.9 Interpret As SInt

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
asSInt	(e)	()	(UInt)	SInt	w_e
			(SInt)	SInt	w_e

The interpret as SInt operation reinterprets e 's bits as a signed integer according to two's complement representation.

7.10 Shift Left Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
shl	(e)	(n)	(UInt) (SInt)	UInt SInt	$w_e + n$ $w_e + n$

The shift left operation concatenates n zero bits to the least significant end of e .

7.11 Shift Right Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
shr	(e)	(n)	(UInt) (SInt)	UInt SInt	$w_e - n$ $w_e - n$

The shift right operation truncates the least significant n bits from e . An error is thrown if e 's bit width is smaller than or equal to n .

7.12 Dynamic Shift Left Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
dshl	(e, n)	()	(UInt, UInt) (SInt, UInt)	UInt SInt	w_e w_e

The dynamic shift left operation shifts the bits in e n places towards the most significant bit. n zeroes are shifted in to the least significant bits, and the n most significant bits are truncated.

7.13 Dynamic Shift Right Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
dshr	(e, n)	()	(UInt, UInt) (SInt, UInt)	UInt SInt	w_e w_e

The dynamic shift right operation shifts the bits in e n places towards the least significant bit. n signed or zeroed bits are shifted in to the most significant bits, and the n least significant bits are truncated.

7.14 Arithmetic Convert to Signed Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
cvt	(e)	()	(UInt) (SInt)	SInt SInt	$w_e + 1$ w_e

The result of the arithmetic convert to signed operation is a signed integer representing the same numerical value as e .

7.15 Negate Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
neg	(e)	()	(UInt) (SInt)	SInt SInt	$w_e + 1$ $w_e + 1$

The result of the negate operation is a signed integer representing the negated numerical value of e .

7.16 Bitwise Complement Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
not	(e)	()	(UInt) (SInt)	UInt UInt	w_e w_e

The bitwise complement operation performs a logical not on each bit in e .

7.17 Binary Bitwise Operations

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
and, or, xor	(e1, e2)	()	(UInt, UInt)	UInt	$\max(w_{e1}, w_{e2})$
			(UInt, SInt)	UInt	$\max(w_{e1}, w_{e2})$
			(SInt, UInt)	UInt	$\max(w_{e1}, w_{e2})$
			(SInt, SInt)	UInt	$\max(w_{e1}, w_{e2})$

The above bitwise operations perform a bitwise and, or, and exclusive or on **e1** and **e2**. The result has the same width as its widest argument, and any narrower arguments are automatically zero-extended or sign-extended to match the width of the result before performing the operation.

7.18 Bitwise Reduction Operations

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
andr, orr, xorr	(e)	()	(UInt)	UInt	1
			(SInt)	UInt	1

The bitwise reduction operations corresponds to a bitwise and, or, and exclusive or operation respectively, reduced over every bit in **e**.

7.19 Concatenate Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
cat	(e1, e2)	()	(UInt, UInt)	UInt	$w_{e1} + w_{e2}$
			(UInt, SInt)	UInt	$w_{e1} + w_{e2}$
			(SInt, UInt)	UInt	$w_{e1} + w_{e2}$
			(SInt, SInt)	UInt	$w_{e1} + w_{e2}$

The result of the concatenate operation is the bits of **e1** concatenated to the most significant end of the bits of **e2**.

7.20 Bit Extraction Operation

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
bits	(e)	(hi, lo)	(UInt)	UInt	hi - lo + 1
			(SInt)	UInt	hi - lo + 1

The result of the bit extraction operation are the bits of **e** between **lo** (inclusive) and **hi** (inclusive). **hi** must be greater than or equal to **lo**. Both **hi** and **lo** must be non-negative and strictly less than the bit width of **e**.

7.21 Head

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
head	(e)	(n)	(UInt)	UInt	n
			(SInt)	UInt	n

The result of the head operation are the **n** most significant bits of **e**. **n** must be positive and less than or equal to the bit width of **e**.

7.22 Tail

Name	Arguments	Parameters	Arg Types	Result Type	Result Width
tail	(e)	(n)	(UInt)	UInt	$w_e - n$
			(SInt)	UInt	$w_e - n$

The tail operation truncates the **n** most significant bits from **e**. **n** must be non-negative and strictly less than the bit width of **e**.

8 Genders

An expression's gender partially determine the legality of connecting to and from the expression. Every expression is classified as either male, female, or bi-gender. For details on connection rules refer back to section ??.

The gender of a reference to a declared circuit component depends on the kind of circuit component. A reference to an input port, an instance, a

memory, and a node, is male. A reference to an output port is female. A reference to a wire or register is bi-gender.

The gender of a subindex or subaccess expression is the gender of the vector-typed expression it indexes or accesses.

The gender of a subfield expression depends upon the orientation of the field. If the field is not flipped, its gender is the gender of the bundle-typed expression it selects its field from. If the field is flipped, then its gender is the reverse of the gender of the bundle-typed expression it selects its field from. The reverse of male is female, and vice-versa. The reverse of bi-gender remains bi-gender.

The gender of all other expressions are male.

9 Namespaces

All modules in a circuit exist in the same module namespace, and thus must all have a unique name.

Each module has an identifier namespace containing the names of all port and circuit component declarations. Thus, all declarations within a module must have unique names. Furthermore, the set of component declarations within a module must be *prefix unique*. Please see section ?? for the definition of prefix uniqueness.

Within a bundle type declaration, all field names must be unique.

During the lowering transformation, all circuit component declarations with aggregate types are rewritten as a sequence of component declarations, each with a ground type. The expansion algorithm in section ?? calculates the names of all replacement components derived from the original aggregate-typed component.

9.1 Expansion Algorithm

Given a component with a ground type, the name of the component is returned.

Given a component with a vector type, the suffix “\$i” is appended to the expanded names of each subelement, where i is the index of each subelement.

Given a component with a bundle type, the suffix “\$f” is appended to the expanded names of each subelement, where f is the field name of each subelement.

9.2 Prefix Uniqueness

A set of names are defined to be *prefix unique* if no name shares a prefix with any other name.

The symbol sequence of a name is the ordered list of strings that results from splitting the name at each occurrence of the ‘\$’ character.

Two names share a prefix if all strings in one’s symbol sequence match exactly to the corresponding strings in the other’s symbol sequence.

As an example `firetruckyz` shares a prefix with `firetruck$y` and `firetruck`, but does not share a prefix with `fire`.

10 Details about Syntax

FIRRTL’s syntax is designed to be human-readable but easily parsed.

The following characters are allowed in identifiers: upper and lower case letters, digits, as well as the punctuation characters `~!@#$$%^*-_+=?/`. Identifiers cannot begin with a digit.

An integer literal in FIRRTL begins with either a hyphen or a digit, and is followed by only digits.

Comments begin with a semicolon and extend until the end of the line. Commas are treated as whitespace, and may be used by the user for clarity if desired.

Block structuring is indicated using indentation. Statements are combined into statement groups by surrounding them with parenthesis. A colon at the *end of a line* will automatically surround the next indented region with parenthesis and thus create a statement group.

The following statement:

```
when c :
  a <= b
else :
  c <= d
  e <= f
```

can be equivalently expressed on a single line as follows.

```
when c : (a <= b) else : (c <= d, e <= f)
```

All circuits, modules, ports and statements can optionally be preceded with the info token `@["filename", line, col]` to annotate them with the source file information from where they were generated.

The following example shows the info tokens included:

```
@["myfile.txt" 14, 8] circuit Top :  
  @["myfile.txt" 15, 2] module Top :  
    @["myfile.txt" 16, 3] output out:UInt  
    @["myfile.txt" 17, 3] input b:UInt<32>  
    @["myfile.txt" 18, 3] input c:UInt<1>  
    @["myfile.txt" 19, 3] input d:UInt<16>  
    @["myfile.txt" 21, 8] wire a:UInt  
    @["myfile.txt" 24, 8] when c :  
      @["myfile.txt" 27, 16] a <= b  
    else :  
      @["myfile.txt" 29, 17] a <= d  
    @["myfile.txt" 34, 4] out <= add(a,a)
```

11 The Lowered FIRRTL Form

The lowered FIRRTL form, LoFIRRTL, is a minimal subset of the FIRRTL language that removes many of the more complicated features. Any FIRRTL compiler must provide a *lowering transformation* that transforms an arbitrary FIRRTL circuit into an equivalent LoFIRRTL circuit.

LoFIRRTL consists of a FIRRTL circuit with the following additional restrictions:

- No aggregate types
- No partial connect statements
- No conditional statements
- All widths have been inferred
- All components are connected to strictly once

This a minimally featured subset is convenient for simple low-level transforms that do not need to support the full FIRRTL feature set.

12 TODO

- Make the syntax for the coding examples consistent.
- Are module names capitalized?
- Is there a space between declared components and the colon before their types.
- etc..
- Make sure all examples are actually correct in every way (all components are initialized)
- For primitive operations make sure we specify what constitutes a legal parameter. (positive?)
- Add short notation section in the beginning of primop section for defining we.
- For tables for primops be consistent with spacing.
- FIRRTL implementation
- Make register reset/init optional
- Rework readwrite port types
- Add memory read-under-write flag
- Add partial connect algorithm
- Add oriented types to type checker
- Add is invalid
- Add validif
- Add UBits
- Add SBits
- Add Mux expression

13 FIRRTL Language Definition

13.1 Abstract Syntax Tree

<i>circuit</i>	=	circuit id : (<i>module</i> [*])	Circuit
<i>module</i>	=	[<i>info</i>] module id : (<i>port</i> [*] <i>stmt</i>)	Module
		[<i>info</i>] extmodule id : (<i>port</i> [*])	External Module
<i>port</i>	=	[<i>info</i>] <i>dir</i> id : <i>type</i>	Port
<i>dir</i>	=	input output	Port Direction
<i>type</i>	=	UInt < <i>width</i> >	Unsigned Integer
		SInt < <i>width</i> >	Signed Integer
		Clock	Clock
		{ <i>field</i> [*] }	Bundle
		<i>type</i> [<i>int</i>]	Vector
<i>field</i>	=	<i>orientation</i> id : <i>type</i>	Bundle Field
<i>orientation</i>	=	default reverse	Orientation
<i>width</i>	=	int ?	Known/Unknown Integer Width
<i>stmt</i>	=	[<i>info</i>] wire id : <i>type</i>	Wire Declaration
		[<i>info</i>] reg id : <i>type</i> , <i>exp</i> , <i>exp</i> , <i>exp</i>	Register Declaration
		[<i>info</i>] mem id : <i>type</i> , int, int, int, (id [*]), (id [*]), (id [*])	Memory Declaration
		[<i>info</i>] inst id : id	Instance Declaration
		[<i>info</i>] poison id : <i>type</i>	Poison Declaration
		[<i>info</i>] node id = <i>exp</i>	Node Declaration
		[<i>info</i>] <i>exp</i> <= <i>exp</i>	Connect
		[<i>info</i>] <i>exp</i> <- <i>exp</i>	Partial Connect
		[<i>info</i>] when <i>exp</i> : <i>stmt</i> else : <i>stmt</i>	Conditional
		[<i>info</i>] stop (<i>exp</i> , <i>exp</i> , int)	Stop Statement
		[<i>info</i>] printf (<i>exp</i> , <i>exp</i> , string, <i>exp</i> [*])	Printf Statement
		[<i>info</i>] skip	Empty Statement
		[<i>info</i>] (<i>stmt</i> [*])	Statement Group
<i>exp</i>	=	[<i>info</i>] UInt < <i>width</i> >(int)	Literal Unsigned Integer
		[<i>info</i>] SInt < <i>width</i> >(int)	Literal Signed Integer
		[<i>info</i>] id	Reference
		[<i>info</i>] <i>exp</i> .id	Subfield
		[<i>info</i>] <i>exp</i> [int]	Subindex
		[<i>info</i>] <i>exp</i> [<i>exp</i>]	Subaccess
		[<i>info</i>] <i>primop</i> (<i>exp</i> [*] , int [*])	Primitive Operation
<i>info</i>	=	filename : line . col	File Location
		noinfo	No File Location

<i>primop</i>	=	add	Unsigned/Signed Add
		sub	Unsigned/Signed Subtract
		addw	Unsigned/Signed Add Wrap
		subw	Unsigned/Signed Subtract Wrap
		mul	Unsigned/Signed Multiply
		div	Unsigned/Signed Divide
		mod	Unsigned/Signed Modulo
		quo	Unsigned/Signed Quotient
		rem	Unsigned/Signed Remainder
		lt	Unsigned/Signed Less Than
		leq	Unsigned/Signed Less or Equal
		gt	Unsigned/Signed Greater Than
		geq	Unsigned/Signed Greater or Equal
		eq	Unsigned/Signed Equal
		neq	Unsigned/Signed Not-Equal
		mux	Unsigned/Signed/Clock Multiplex
		pad	Unsigned/Signed Pad to Length
		asUInt	Unsigned/Signed Reinterpret Bits as UInt
		asSInt	Unsigned/Signed Reinterpret Bits as SInt
		shl	Unsigned/Signed Shift Left
		shr	Unsigned/Signed Shift Right
		dshl	Unsigned/Signed Dynamic Shift Left
		dshr	Unsigned/Signed Dynamic Shift Right
		cvt	Unsigned/Signed to Signed Logical Conversion
		neg	Unsigned/Signed Negate
		not	Unsigned Not
		and	Unsigned And
		or	Unsigned Or
		xor	Unsigned Xor
		andr	Unsigned And Reduce
		orr	Unsigned Or Reduce
		xorr	Unsigned Xor Reduce
		cat	Unsigned Concatenation
		bit	Single Bit Extraction
		bits	Multiple Bit Extraction
		toClock	Interpret Unsigned Bit as Clock
		fromClock	Interpret Clock as Unsigned Bit

13.2 Notation

The above definition specifies the structure of the abstract syntax tree corresponding to a FIRRTL circuit. Nodes in the abstract syntax tree are

italicized. Keywords are shown in **bold**. The special productions `id`, `int`, and `string`, indicates an identifier, an integer literal, and a string respectively. Tokens followed by an asterisk, *e.g.* `field*`, indicates a list formed from repeated occurrences of the token.

Keep in the mind that the above definition is only the *abstract* syntax tree, and is a representation of the in-memory FIRRTL data structure. Readers and writers are provided for converting a FIRRTL data structure into a purely textual representation, which is defined in Section ??.